

```
import warnings
warnings.filterwarnings("ignore")
import numpy as np
from numpy import array
import pandas as pd
from tqdm import tqdm
import datetime as datetime
import tensorflow as tf
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from wordcloud import WordCloud
import glob
import string
import cv2
import random
import imgaug.augmenters as iaa
from pickle import dump, load
from keras.preprocessing import image
from keras.preprocessing.image import load_img, img_to_array
from keras.applications.vgg16 import VGG16, preprocess_input
from keras.layers import LSTM, Embedding, Dense, Activation, Dropout
from keras.callbacks import ModelCheckpoint, ModelCheckpoint, TensorBoard
from keras.layers.merge import Add
from keras.preprocessing import sequence
from keras.models import Model
from keras import Input, layers
from keras import optimizers
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from nltk.translate.bleu_score import sentence_bleu
tf.random.set_seed(3)

image_path = '/content/Flicker8k_Dataset/'
dir_Flickr_text = '/content/Flickr_txt/Flickr8k.token.txt'
images_lst = glob.glob(image_path + '*.jpg')
no_of_image = len(images_lst)
print(f'Total Images in Datset : {no_of_image}')

Total Images in Datset : 8091
```

Utility Functions:

- Load text
- Preparing Text data from Flickr8k.txt file
- Prepare DataFrames

```

def load_txt(txt_file):
    with open(txt_file, 'r') as file:
        text = file.read()
    return text

def prepare_clean_txt(text):
    text_data = []
    table = str.maketrans('', '', string.punctuation)
    for line in tqdm(text.split('\n')):
        col = line.split('\t')
        if len(col) == 1:
            continue
        filename, index = col[0].split("#")
        filename_path = image_path + filename
        txt = col[1].split()
        txt = [word.lower() for word in txt]
        txt = [w.translate(table) for w in txt]
        txt = [word for word in txt if len(word)>1]
        txt = [word for word in txt if word.isalpha()]
        txt = ' '.join(txt)
        text_data.append([index] + [filename_path] + [txt])
    return text_data

```

splitting over each line
splitting over each space in
neglecting data if full info

```

def prepare_dataframes(text_data):
    data_df = pd.DataFrame(text_data, columns=[ 'index', 'filename', 'caption'])
    data_df = data_df[data_df.filename != '2258277193_586949ec62.jpg.1']
    data_df['caption_with_tags'] = data_df['caption'].apply(lambda txt : 'startseq ' + txt + ' ')
    tokenizer = Tokenizer()
    tokenizer.fit_on_texts(data_df.caption.values)
    data_df['caption_word_count'] = data_df['caption'].apply(lambda caption: len(caption.split()))
    word = []
    frequency = []
    for k,v in tokenizer.word_counts.items():
        word.append(k)
        frequency.append(v)
    df_word_count = pd.DataFrame({'word': word, 'frequency': frequency})
    df_word_count.sort_values('frequency', ascending=False, inplace=True, ignore_index=True)
    return data_df, df_word_count

```

splitting to gather information
convert each word to lower case
to remove punctuations from
to discard words of length 1
to discard numeric string

```

text = load_txt(dir_Flickr_text)
text_data = prepare_clean_txt(text)
data_df, df_word_count = prepare_dataframes(text_data)

```

100% |██████████| 40461/40461 [00:00<00:00, 64152.52it/s]

data_df.head(3)

▼ Exploratory Data Analysis

```
from collections import Counter
def utility_counter(data):
    filename = data_df.filename.values
    caption_count = Counter(filename).values()
    uni_filenames = np.unique(filename)
    print(f'Number of unique file names : {len(uni_filenames)}')
    print(f'Number of captions per image : {set(caption_count)}')
    print(f'All image have same number of captions: {all(caption_count)}')

utility_counter(data_df)

Number of unique file names : 8092
Number of captions per image : {5}
All image have same number of captions: True
```

Image - Caption visualization

```
def image_caption_plotter(data):
    npic = 3
    target_size = (350, 500, 3)
    count = 1
    fig = plt.figure(figsize=(25, 22))
    for file in tqdm(random.sample(list(data['filename']), npic)):
        captions = list(data['caption'].loc[data['filename'] == file].values)
        image_load = load_img(file, target_size = target_size)
        ax = fig.add_subplot(npic, 2, count)
        ax.axis('off')
        ax.imshow(image_load, interpolation='nearest')
        count += 1
        ax = fig.add_subplot(npic, 2, count)
        ax.axis('off')
        ax.set_ylim(0, len(captions))
        for i, caption in enumerate(captions):
            ax.text(0, i, caption, fontsize = 30, fontfamily= 'fantasy', bbox=dict(facecolor='black'))
        count += 1
        plt.grid(None)
    plt.show()

image_caption_plotter(data_df)
```

100% | [3/3 [00:00<00:00, 29.35it/s]



THE MAN IS INTERVIEWING TEENAGE BOY IN FRONT OF AN AUDIENCE
MAN INTERVIEWS BOY IN FRONT OF CROWD OF PEOPLE
MAN INTERVIEWS BOY IN GATHERING OF PEOPLE
GROUP OF PEOPLE WATCHING BOY GETTING INTERVIEWED BY MAN
CHILD BEING INTERVIEWED BY MAN WITH MICROPHONE



DOG LOOKING AT TENNIS BALL FLOATING IN POOL
DOG DESCENDS RAMP INTO THE POOL TO GET TENNIS BALL
BROWN AND WHITE DOG WALKS DOWN RAMP TOWARD POOL OF WATER IN WHICH GREEN TENNIS BALL FLOATS
BROWN AND WHITE DOG LOOKING INTO POOL AT TENNIS BALL
BROWN AND WHITE DOG IS LOOKING AT TENNIS BALL FLOATING IN POOL



THE GIRL DRESSED IN PINK AND WHITE AND WEARING KNIT CAP IS REACHING TOWARD FLOWERING TREE
YOUNG GIRL WITH WHITE VEST PINK SLEEVES AND PINK KNIT HAT WITH FLOWER IS LOOKING AT THE FLOWER BLOSSOMS
YOUNG GIRL IS PICKING BLOSSOMS FROM TREE
GIRL IN PINK HAT PICK FLOWER FROM TREE
CHILD IN WHITE AND PINK IS PICKING WHITE FLOWERED TREE NEXT TO FENCE

Observation

- We can see each image has 5 captions associated with it.

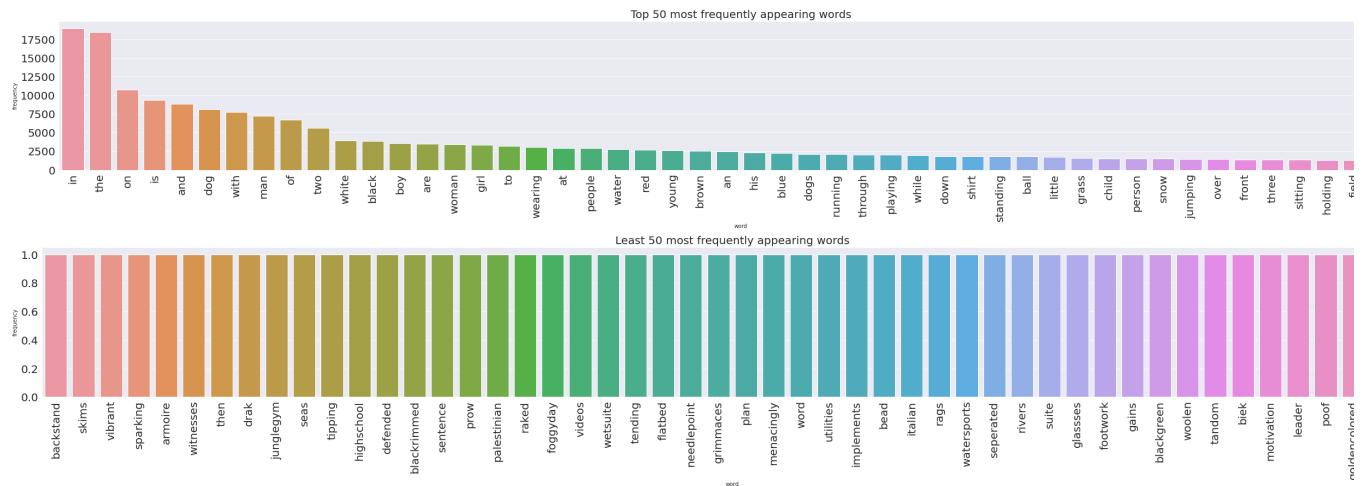
Plot for word - frequency

count = 50

```
def plot_word_count(df, title = ''):
    sns.set_style('darkgrid')
    plt.figure(figsize = (45, 5))
    sns.barplot(x=df['word'], y=df["frequency"])
```

```
plt.yticks(fontsize = 20)
plt.xticks( rotation = 'vertical', fontsize = 20)
plt.title(title, fontsize = 20)
plt.show()
```

```
plot_word_count(df_word_count.iloc[:count, :], title= 'Top 50 most frequently appearing words  
plot_word_count(df_word_count.iloc[-count:, :], title= 'Least 50 most frequently appearing wo
```



Word Cloud Visualization

```
text = ''  
for caption in data_df.caption:  
    text = text + caption  
plt.figure(figsize=(16, 12))  
word_cloud = WordCloud(width = 1000, height = 500, random_state=1, collocations= True, backgr  
plt.imshow(word_cloud, interpolation='bilinear')  
plt.axis("off")  
plt.title('Strength of word in complete data corpus', fontsize = 30)  
plt.show()
```

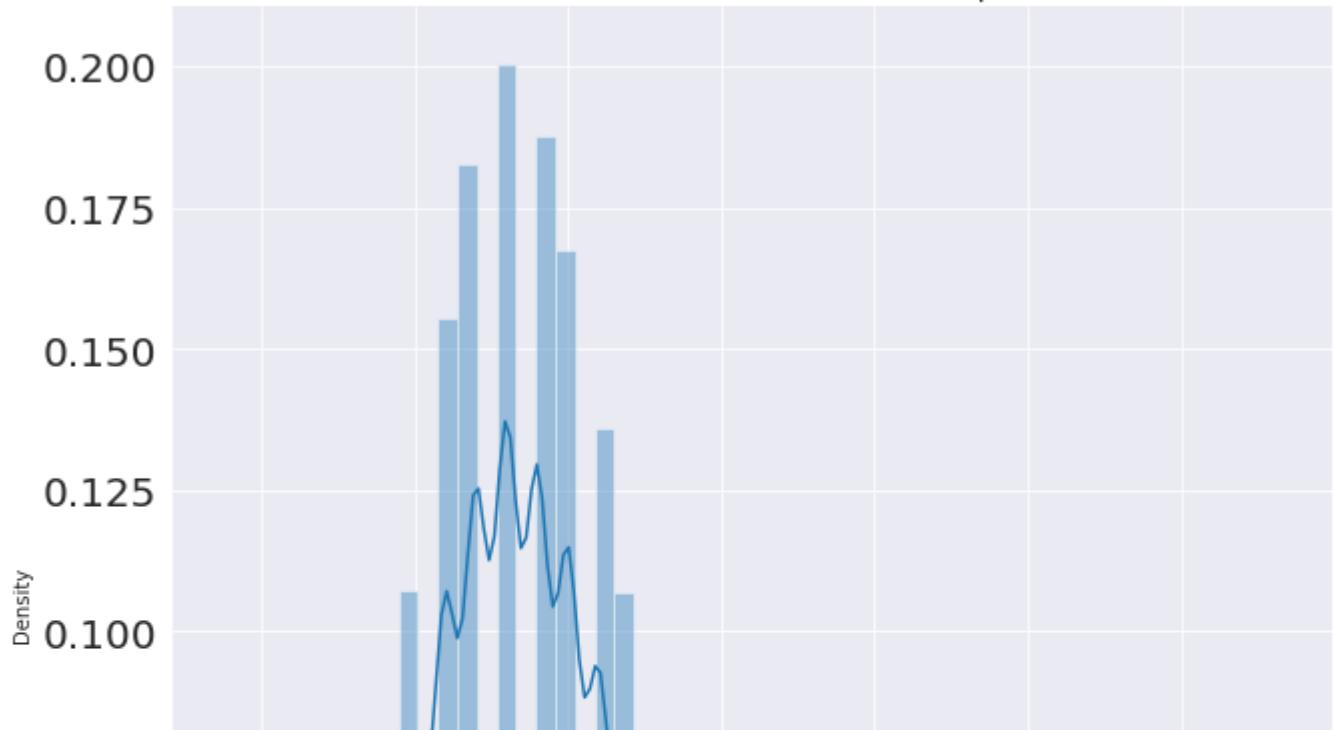
Strength of word in complete data corpus



Distribution of number of words in each caption

```
sns.set_style('darkgrid')
sns.FacetGrid(data_df, height = 9).map(sns.distplot, 'caption_word_count').add_legend()
plt.yticks(fontsize = 20)
plt.xticks(fontsize = 20)
plt.title('Word Count distribution in each Caption', fontsize = 15)
plt.show()
```

Word Count distribution in each Caption



Observation:

- We can see most of captions contain words in range of 7 to 10

- Train_filename_container : Train image path
- Test_filename_container : Test image path

```

train_images_file = '/content/Flickr_txt/Flickr_8k.trainImages.txt'
train_images = set(open(train_images_file, 'r').read().strip().split('\n'))
train_filename_container = [image_path+filename for filename in train_images if image_path+fi

validation_images_file = '/content/Flickr_txt/Flickr_8k.devImages.txt'
validation_images = set(open(validation_images_file, 'r').read().strip().split('\n'))
validation_filename_container = [image_path+filename for filename in validation_images if ima

train_filename_container = train_filename_container+validation_filename_container

test_images_file = '/content/Flickr_txt/Flickr_8k.testImages.txt'
test_images = set(open(test_images_file, 'r').read().strip().split('\n'))
test_filename_container = [image_path+filename for filename in test_images if image_path+file

```

Augmentation on images

```

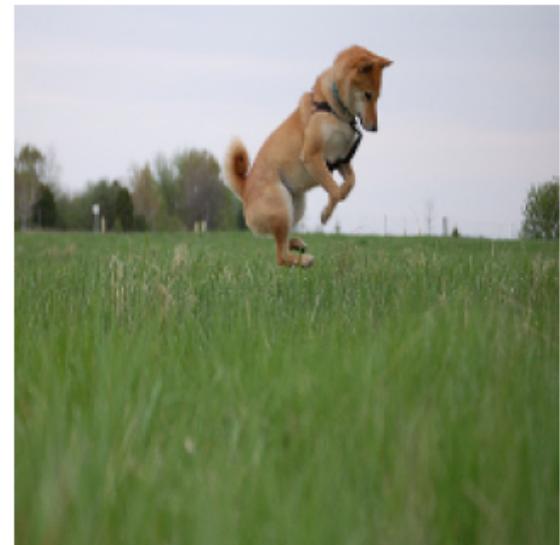
def augment_image(image_path, aug=True, visualize=False):
    if visualize:
        image = cv2.imread(image_path, cv2.IMREAD_UNCHANGED)

```

```
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
image = cv2.resize(image, (224, 224))
else:
    image = load_img(image_path, target_size=(224, 224))
    image = img_to_array(image)
if aug:
    aug2 = iaa.Fliplr(1)
    aug3 = iaa.Flipud(1)
    aug4 = iaa.Emboss(alpha=(1), strength=1)
    aug5 = iaa.DirectedEdgeDetect(alpha=(0.8), direction=(1.0))
    aug6 = iaa.Sharpen(alpha=(1.0), lightness=(1.5))

    a = np.random.uniform()
    if a<0.2:
        aug_image = aug2.augment_image(image)
    elif a<0.4:
        aug_image = aug3.augment_image(image)
    elif a<0.6:
        aug_image = aug4.augment_image(image)
    elif a<0.8:
        aug_image = aug5.augment_image(image)
    else:
        aug_image = aug6.augment_image(image)
return image, aug_image
else:
    return image
```

```
row = 3
col = 2
dataset = [augment_image(train_filename_container[np.random.randint(0, 5000)], aug=True, vis=False) for _ in range(5000)]
fig = plt.figure(figsize=(15, 20))
for counter in range(0, row*col):
    fig.add_subplot(row, col, counter+1)
    plt.imshow((dataset[int(counter/2)][counter%2]))
    plt.axis('off')
    plt.grid(None)
```



- train_caption_mapping: Train image caption
 - test_caption_mapping: Train image caption
-

```
train_caption_mapping = {}
```

```

for file in tqdm(train_filename_container):
    train_caption_mapping[file] = list(data_df.loc[data_df['filename'] == file].caption_with_t

test_caption_mapping = {}
for file in tqdm(test_filename_container):
    test_caption_mapping[file] = list(data_df.loc[data_df['filename'] == file].caption_with_ta

```

100%|██████████| 7000/7000 [00:20<00:00, 336.60it/s]
100%|██████████| 1000/1000 [00:02<00:00, 374.27it/s]

Model load VGG16

```

def encode_image(filename_container, aug=True):
    encoding = {}
    augmented_image_path = []
    for image_path in tqdm(filename_container):
        image = augment_image(image_path, aug=aug)
        if aug:
            image = preprocess_input(image[0])
            aug_image = preprocess_input(image[1])

            features = model_new(np.expand_dims(image, axis=0))
            features = np.reshape(features, (features.shape[1],))
            features_aug = model_new(np.expand_dims(image, axis=0))
            features_aug = np.reshape(features_aug, (features_aug.shape[1],))

            encoding[image_path] = features
            aug_path = image_path[:-4] + '_aug.jpg'
            augmented_image_path.append(aug_path)
            train_caption_mapping[aug_path] = train_caption_mapping[image_path]
            encoding[aug_path] = features_aug
        else:
            image = preprocess_input(image)
            features = model_new(np.expand_dims(image, axis=0))
            features = np.reshape(features, (features.shape[1],))
            encoding[image_path] = features
    global train_filename_container
    train_filename_container += augmented_image_path
    return encoding

```

```

model_vgg16 = VGG16(weights='imagenet', include_top=True)
model_new = Model(model_vgg16.input, model_vgg16.layers[-2].output)

```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/vgg16/>
553467904/553467096 [=====] - 10s 0us/step
553476096/553467096 [=====] - 10s 0us/step



Save the bottleneck train and test features to disk

```
encoding_train = encode_image(train_filename_container)
with open('/content/encoded_train_image_features.pkl', 'wb') as p:
    dump(encoding_train, p)
```

100%|██████████| 7000/7000 [03:18<00:00, 35.26it/s]

```
encoding_test = encode_image(test_filename_container, aug=False)
with open('/content/encoded_test_image_feature.pkl', 'wb') as p:
    dump(encoding_test, p)
```

100%|██████████| 1000/1000 [00:19<00:00, 50.46it/s]

Loading saved features

```
train_features = load(open("/content/encoded_train_image_features.pkl", "rb"))
print('Number of train image features', len(train_features))

test_features = load(open("/content/encoded_test_image_feature.pkl", "rb"))
print('Number of test image features', len(test_features))
```

Number of train image features 14000

Number of test image features 1000

List of all the train, test captions

```
all_trainCaptions = set()
for _, caption_lst in tqdm(trainCaptionMapping.items()):
    for caption in caption_lst:
        all_trainCaptions.add(caption)

allTestCaptions = set()
for _, caption_lst in tqdm(testCaptionMapping.items()):
    for caption in caption_lst:
        allTestCaptions.add(caption)
```

100%|██████████| 14000/14000 [00:00<00:00, 520994.57it/s]

100%|██████████| 1000/1000 [00:00<00:00, 498195.04it/s]

Frequency of words in train, test caption

```
trainTokenizer = Tokenizer(oov_token = '<UNK>')
trainTokenizer.fit_on_texts(all_trainCaptions)
words = []
frequency = []
```

```

for word, freq in tqdm(train_tokenizer.word_counts.items()):
    words.append(word)
    frequency.append(freq)
df_wc_train = pd.DataFrame({'word': words, 'frequency': frequency})
df_wc_train.sort_values('frequency', ascending=False, inplace=True, ignore_index=True)

100%|██████████| 8169/8169 [00:00<00:00, 1144130.28it/s]

```

Word-Index/Index-Word Mapping

```

ixtoword = dict((index, word) for word, index in train_tokenizer.word_index.items())
wordtoix = train_tokenizer.word_index
max_length = max(map(lambda caption: len(caption.split()), all_train_captions))
vocab_size = len(wordtoix)+1

```

Dataset Preparation/Dataloader

```

def dataset_loader(tokenizer, image_caption_mapping, image_features, max_length, images_per_batch):
    image_feature, input_seq, output_seq = list(), list(), list()
    image_count=0
    while 1:
        for image_path, captions in image_caption_mapping.items():
            image_count+=1
            feature = image_features[image_path]
            sequences = tokenizer.texts_to_sequences(captions)
            for seq in sequences:
                for i in range(1, len(seq)):
                    in_seq, out_seq = seq[:i], seq[i]
                    in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
                    out_seq = tf.keras.utils.to_categorical([out_seq], num_classes=vocab_size)[0]
                    image_feature.append(feature)
                    input_seq.append(in_seq)
                    output_seq.append(out_seq)
            if image_count==images_per_batch:
                yield [[array(image_feature), array(input_seq)], array(output_seq)]
                image_feature, input_seq, output_seq = list(), list(), list()
                image_count=0

```

Embeddings: Load Glove vectors

```

embeddings_index = {}
f = open('/content/glove.6B.200d.txt', encoding="utf-8")

for line in tqdm(f):
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')

```

```

embeddings_index[word] = coefs
f.close()
print(f'\n Number of words in glove vector: {len(embeddings_index)}')

400000it [00:14, 27148.35it/s]
Number of words in glove vector: 400000

```

Embedding matrix

```

embedding_dim = 200
embedding_matrix = np.zeros((vocab_size, embedding_dim))
for word, i in wordtoix.items():
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        embedding_matrix[i] = embedding_vector

embedding_matrix.shape

(8171, 200)

```

Model Architecture

```

input_image = Input(shape=(4096, ), name='Image_Feature_input')
fe1 = Dropout(0.5, name='Dropout_image')(input_image)
fe2 = Dense(256, activation='relu', name='Activation_Encoder')(fe1)
input_text = Input(shape=(max_length,), name='Text_input')
se1 = Embedding(vocab_size, embedding_dim, mask_zero=True, weights=[embedding_matrix], traina
se2 = Dropout(0.5, name='Dropout_text')(se1)
se3 = LSTM(256, name='LSTM')(se2)
decoder1 = Add(name='Add')([fe2, se3])
decoder2 = Dense(256, activation='relu', name='Activation_Decoder')(decoder1)
output = Dense(vocab_size, activation='softmax', name='Output')(decoder2)
model = Model(inputs=[input_image, input_text], outputs=output)

```

Model Summary

```
model.summary()
```

```
Model: "model_1"
```

| Layer (type) | Output Shape | Param # | Connected to |
|----------------------------------|----------------|---------|--------------|
| <hr/> | | | |
| Text_input (InputLayer) | [(None, 34)] | 0 | [] |
| Image_Feature_input (InputLayer) | [(None, 4096)] | 0 | [] |

| | | | |
|----------------------------|-----------------|---------|--|
| Text_Feature (Embedding) | (None, 34, 200) | 1634200 | ['Text_input[0][0]'] |
| Dropout_image (Dropout) | (None, 4096) | 0 | ['Image_Feature_input[0]'] |
| Dropout_text (Dropout) | (None, 34, 200) | 0 | ['Text_Feature[0][0]'] |
| Activation_Encoder (Dense) | (None, 256) | 1048832 | ['Dropout_image[0][0]'] |
| LSTM (LSTM) | (None, 256) | 467968 | ['Dropout_text[0][0]'] |
| Add (Add) | (None, 256) | 0 | ['Activation_Encoder[0][0]', 'LSTM[0][0]'] |
| Activation_Decoder (Dense) | (None, 256) | 65792 | ['Add[0][0]'] |
| Output (Dense) | (None, 8171) | 2099947 | ['Activation_Decoder[0][0]'] |

=====

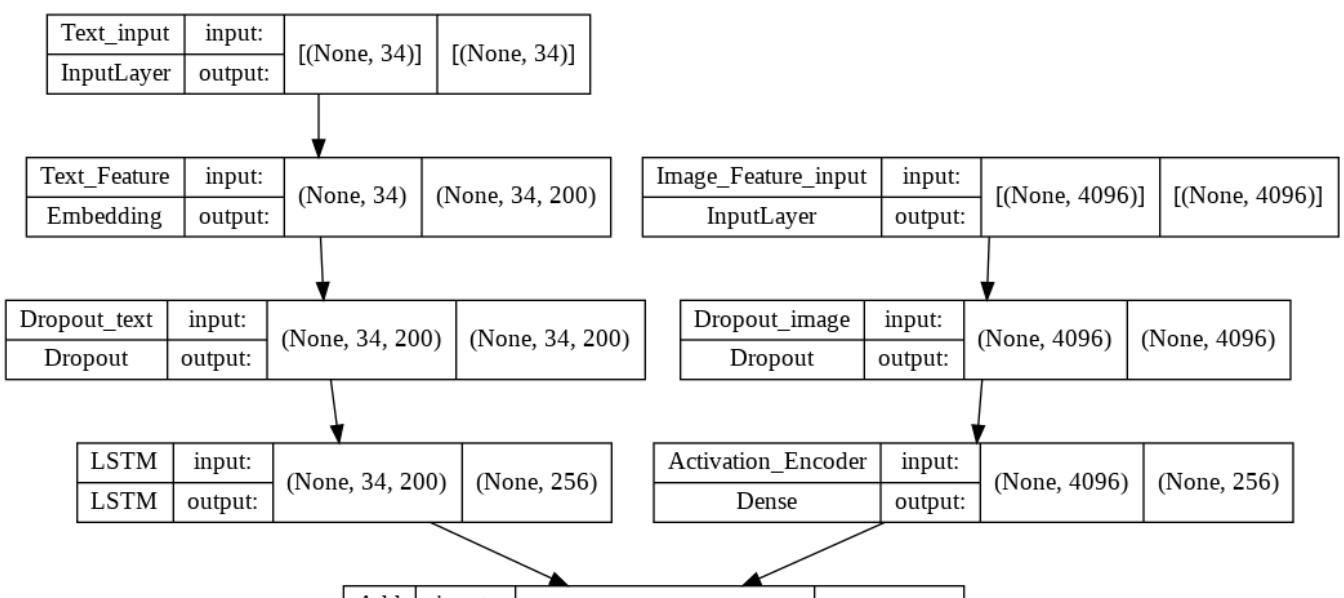
Total params: 5,316,739

Trainable params: 3,682,539

Non-trainable params: 1,634,200

Model Architecture

```
tf.keras.utils.plot_model(model, to_file='model.png', show_shapes = True, show_layer_names =
```



Compile/Callbacks

```

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
filepath = 'model-ep{epoch:03d}-loss{loss:.3f}.h5'
checkpoint = ModelCheckpoint(filepath, monitor='loss', verbose=1, save_best_only=True, mode='log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")'
tensorboard = TensorBoard(log_dir = log_dir, histogram_freq=1)
  
```

Model Fitting

```

number_pics_per_bath = 16
steps = len(train_caption_mapping)//number_pics_per_bath
learning_rate = 0.001

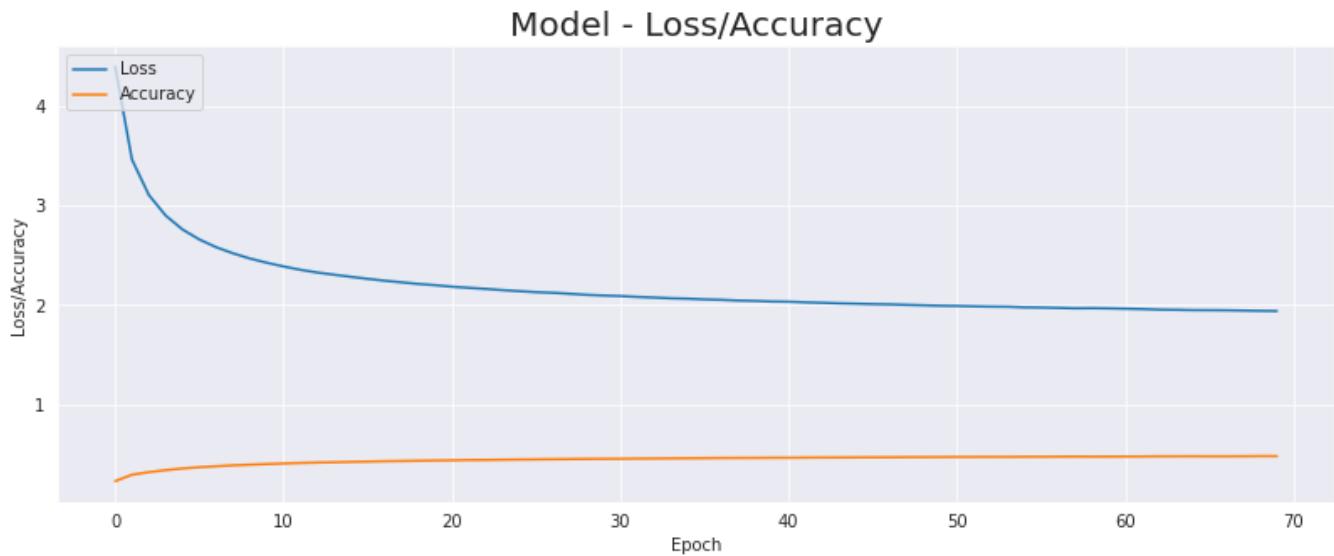
train_generator = dataset_loader(train_tokenizer, train_caption_mapping, train_features, max_
history = model.fit(train_generator, epochs=70, steps_per_epoch=steps, verbose=1, callbacks =
    ...,
Epoch 57/70
875/875 [=====] - ETA: 0s - loss: 1.9713 - accuracy: 0.4743
Epoch 57: loss improved from 1.97496 to 1.97126, saving model to model-ep057-loss1.97
875/875 [=====] - 151s 173ms/step - loss: 1.9713 - accuracy:
Epoch 58/70
875/875 [=====] - ETA: 0s - loss: 1.9672 - accuracy: 0.4755
Epoch 58: loss improved from 1.97126 to 1.96715, saving model to model-ep058-loss1.96
875/875 [=====] - 150s 171ms/step - loss: 1.9672 - accuracy:
Epoch 59/70
875/875 [=====] - ETA: 0s - loss: 1.9683 - accuracy: 0.4745
Epoch 59: loss did not improve from 1.96715
875/875 [=====] - 150s 172ms/step - loss: 1.9683 - accuracy:
Epoch 60/70
875/875 [=====] - ETA: 0s - loss: 1.9661 - accuracy: 0.4750
Epoch 60: loss improved from 1.96715 to 1.96613, saving model to model-ep060-loss1.96
875/875 [=====] - 150s 171ms/step - loss: 1.9661 - accuracy:
Epoch 61/70
  
```

```
875/875 [=====] - ETA: 0s - loss: 1.9631 - accuracy: 0.4755
Epoch 61: loss improved from 1.96613 to 1.96310, saving model to model-ep061-loss1.96
875/875 [=====] - 150s 171ms/step - loss: 1.9631 - accuracy:
Epoch 62/70
875/875 [=====] - ETA: 0s - loss: 1.9596 - accuracy: 0.4760
Epoch 62: loss improved from 1.96310 to 1.95958, saving model to model-ep062-loss1.96
875/875 [=====] - 152s 173ms/step - loss: 1.9596 - accuracy:
Epoch 63/70
875/875 [=====] - ETA: 0s - loss: 1.9536 - accuracy: 0.4775
Epoch 63: loss improved from 1.95958 to 1.95356, saving model to model-ep063-loss1.95
875/875 [=====] - 150s 172ms/step - loss: 1.9536 - accuracy:
Epoch 64/70
875/875 [=====] - ETA: 0s - loss: 1.9519 - accuracy: 0.4778
Epoch 64: loss improved from 1.95356 to 1.95187, saving model to model-ep064-loss1.95
875/875 [=====] - 152s 173ms/step - loss: 1.9519 - accuracy:
Epoch 65/70
875/875 [=====] - ETA: 0s - loss: 1.9488 - accuracy: 0.4782
Epoch 65: loss improved from 1.95187 to 1.94877, saving model to model-ep065-loss1.94
875/875 [=====] - 149s 171ms/step - loss: 1.9488 - accuracy:
Epoch 66/70
875/875 [=====] - ETA: 0s - loss: 1.9481 - accuracy: 0.4777
Epoch 66: loss improved from 1.94877 to 1.94810, saving model to model-ep066-loss1.94
875/875 [=====] - 152s 173ms/step - loss: 1.9481 - accuracy:
Epoch 67/70
875/875 [=====] - ETA: 0s - loss: 1.9472 - accuracy: 0.4776
Epoch 67: loss improved from 1.94810 to 1.94718, saving model to model-ep067-loss1.94
875/875 [=====] - 151s 173ms/step - loss: 1.9472 - accuracy:
Epoch 68/70
875/875 [=====] - ETA: 0s - loss: 1.9443 - accuracy: 0.4784
Epoch 68: loss improved from 1.94718 to 1.94434, saving model to model-ep068-loss1.94
875/875 [=====] - 151s 172ms/step - loss: 1.9443 - accuracy:
Epoch 69/70
875/875 [=====] - ETA: 0s - loss: 1.9414 - accuracy: 0.4796
Epoch 69: loss improved from 1.94434 to 1.94136, saving model to model-ep069-loss1.94
875/875 [=====] - 150s 172ms/step - loss: 1.9414 - accuracy:
Epoch 70/70
875/875 [=====] - ETA: 0s - loss: 1.9394 - accuracy: 0.4792
Epoch 70: loss improved from 1.94136 to 1.93944, saving model to model-ep070-loss1.93
875/875 [=====] - 149s 171ms/step - loss: 1.9394 - accuracy:
```

Loss/Accuracy Plot

```
plt.figure(figsize=(30, 5))
plt.subplot(121)
plt.plot(history.history['loss'])
plt.plot(history.history['accuracy'])
plt.title('Model - Loss/Accuracy', fontsize=20)
plt.ylabel('Loss/Accuracy')
plt.xlabel('Epoch')
plt.legend(['Loss', 'Accuracy'], loc='upper left')
```

```
<matplotlib.legend.Legend at 0x7f37c50d3ed0>
```



Tensorboard Plot

```
%load_ext tensorboard  
%tensorboard --logdir logs/fit
```

TensorBoard

SCALARS

GRAPHS

INACTIVE

- Show data download links
- Ignore outliers in chart scaling

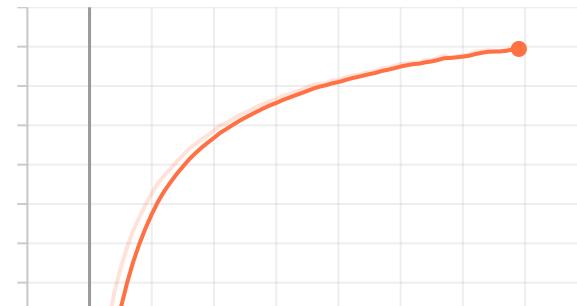
Tooltip sorting method:

Smoothing

 0.6

epoch_accuracy

epoch_accuracy
tag: epoch_accuracy



Inference

ABSOLUTE

```
with open("/content/encoded_test_image_feature.pkl", "rb") as encoded_pickle:
    encoding_test = load(encoded_pickle)
```

epoch loss

```
def greedySearch(image):
    image = encoding_test[image]
    image = np.reshape(image, (1,4096))
    in_text = 'startseq'
    for i in range(max_length):
        sequence = [wordtoix[word] for word in in_text.split()]
        sequence = pad_sequences([sequence], maxlen=max_length)
        yhat = model.predict([image,sequence], verbose=0)
        yhat = np.argmax(yhat)
        word = ixtoword[yhat]
        in_text += ' ' + word
        if word == 'endseq':
            break
    final = in_text.split()
    final = final[1:-1]
    final = ' '.join(final)
    return final
```

```
def image_caption_score_extractor(data):
    npic = 6
    target_size = (224, 224, 3)
    count = 1
    fig = plt.figure(figsize=(60,65))
    for file in tqdm(random.sample(test_filename_container, npic)):
        predicted_caption = greedySearch(file)
        captions = list(data['caption'].loc[data['filename'] == file].values)
        image_load = load_img(file, target_size = target_size)
```

```
image_load = image_load[1::2], caption = caption[1::2])
ax = fig.add_subplot(npic, 2, count)
ax.axis('off')
ax.imshow(image_load, interpolation='nearest')
reference = [caption.split() for caption in captions]
candidate = predicted_caption.split()
score = round(sentence_bleu(reference, candidate), 3)
count += 1
ax = fig.add_subplot(npic, 2, count)
ax.axis('off')
ax.set_ylim(0, len(captions)+2)
for i, caption in enumerate(captions):
    ax.text(0, i, caption, fontsize = 40, bbox=dict(facecolor='blue', alpha=0.1), color='red')
    ax.text(0, i+1, greedySearch(file), fontsize =40, fontfamily= 'fantasy', bbox=dict(facecolor='white', alpha=0.1), color='black')
    ax.text(0, i+2, 'BLEU Score: '+str(score), fontsize = 40, fontfamily= 'fantasy', bbox=dict(facecolor='white', alpha=0.1), color='black')
    count += 1
plt.grid(None)
plt.show()
```

```
image_caption_score_extractor(data_df)
```

100% | 

| 6/6 [00:06<00:00, 1.08s/it]



BLEU SCORE: 0.434

PERSON IN RED JACKET IS CLIMBING STEEP MOUNTAINSIDE COVERED
 skier wearing white sunglasses is skiing down snowy mountain very fast
 skier skiing downhill in cloud of snow
 skier is throwing up snow as he skis off piste
 person dressed in red and grey is skiing down snowy slope
 man is skiing down mountain



BLEU SCORE: 0.481

MAN IN BLACK COAT AND KHAKI EATS CIGARETTE
 the small child bangs his drum while looking at an older man doing the same
 the man with the big drum is marching next to boy with small drum
 man plays drum and little boy hits his own little drum
 man and little boy beating drums
 large man in leather costume plays big drum beside little boy playing little drum



BLEU SCORE: 0.482

BOY HOLDS HIS CELLPHONE TO THE CAMERA
 young child holds spoon to its mouth while sitting in chair
 small child dressed in green is eating with spoon
 little boy is eating his food off of spoon while sitting on patio
 little boy holds spoon up to his mouth
 boy eats with spoon



BLEU SCORE: 0.687

TWO PEOPLE IN RED SHIRTS WALK DOWN THE STREET
 the man is wearing green shirt and blue jeans and carrying free hugs sign
 young man in black shirt is holding free hugs sign
 man is holding sign that says free hugs
 man holds up free hugs sign above his head
 man carrying sign that says free hug along the sidewalk



BLEU SCORE: 0.359

LITTLE GIRL IN PINK SHIRT IS CLIMBING THROUGH THE TUB
 the little girl in the green outfit is playing with piece of orange material
 young girl plays outside
 little girl in green outfit wringing out an orange rag
 girl wearing green twists something in her hands
 girl wrings out wet shirt



BLEU SCORE: 0.788

WOMAN IN COWBOY SUIT AND MAN IN WETSUIT STANDING IN FRONT OF MOUNTAIN
 woman taking pictures of the surf and child in the foreground
 woman is taking photograph on beach at the edge of the water while girl stands in the foreground
 man and girl are at the shoreline while the man takes picture and the girl holds doll
 little girl in red at the beach while adult points camera
 child on the beach while woman takes pictures nearby